

MorphOS


An Extensible Networked Operating System

Peter Okelmann, Ilya Meignan-Masson, Masanori Misono,
Pramod Bhatotia




High-Performance Networking for the Cloud

- Fast, networked applications
 - Serverless apps
 - Virtual Network Functions (VNFs)
- App-OS optimizations and codesign
 - Network stack
 - Scheduling
 - VM instance chaining, ...



OS design is critical for fast, networked applications

OS Architectures for Network Stacks



Unikernels are ideal for cloud VMs and networked applications

Unikernel Characteristics



Pros

- High performance
- Small image size
- Minimized Trusted Computing Base (TCB)

Cons

- Lacking **extensibility**
- Lack of address space isolation
- No multi-tenant safety

Can we get unikernel benefits while mitigating the shortcomings?

Problem Statement




How do we achieve extensibility of fast, networked unikernel applications without compromising safety?

Design Goals

- Reconfigurability
- Safe multi-tenancy
- Performance

Our Proposal: MorphOS

- **eBPF runtime:** safe and fast extension execution
- **Control plane:** reconfigurability through eBPF updates
- **Data plane:** core application logic leveraging MorphOS hookpoints



An extensible networked operating system that brings verified eBPF to unikernels

Outline

- Overview

- Design
- Evaluation

Challenges of Unikernels

Lack of extensibility

Lack of isolation

Solution #1
Hookpoints:
Extensibility with
eBPF


Solution #2
Offline:
Decoupled
verification

Solution #3
Online:
Hardware-assisted
runtime hardening

#1: Extensibility with eBPF Hookpoints

Multi-tenant VM development:

1. Provider builds optimized app
-> **MorphOS Provider API**
2. Tenant build eBPF program
-> **MorphOS Tenant API**
3. JIT compilation
4. Runtime execution



MorphOS APIs make eBPF available to providers and define a runtime for tenants

Challenges of Unikernels

Lack of extensibility

Solution #1
Hookpoints:
extensibility with
eBPF


Lack of isolation

Solution #2
Offline:
Decoupled
verification

Solution #3
Online:
Hardware-assisted
runtime hardening

#2 Decoupled eBPF Verification


- Protect against malicious tenants
 - Integrity, confidentiality, availability
- Verification service
 - Static control flow analysis
 - Track memory values and ranges
 - Generates cryptographic certificate
- Unikernel validates certificate



MorphOS allows providers to dedicate resources for asynchronous verification

#3 MPK: Hardening against Verifier Bugs

- Current eBPF verifiers are error prone
 - Linux verifier is unfit for security purposes
 - Correctness of verifiers remains unproven
⇒ Runtime hardening
- Memory Protection Keys (MPK)
 - 16 protection domains (PKeys)
⇒ Fast domain switches via Wrpkru
- Protection domains: unikernel, eBPF, packets, ...
- Transport packets by changing protection domains



MorphOS hardens eBPF execution even in the presence of verifier/JIT bugs

Outline

- Overview
- Design
- Evaluation

Implementation and Case Study

MorphOS implementation:

- Built on **Unikraft** unikernel (v0.16.3)
- Extend **Prevail** verifier
- Extend the **uBPF** JIT compiler



Case study app: MorphClick

- Click: flexible VNF system
- Replace native algorithms with eBPF hookpoints
 - Firewall: replace ACL tables with JIT compiler
 - DPI (string matching): frequent reconfiguration
 - NAT: retain state across reconfigurations

Evaluation

- **RQ #1:** Does MorphOS improve **live-reconfigurability?**
 - Reconfiguration time
- **RQ #2:** Does MorphOS harden against **verifier bugs?**
 - Effectiveness of hardening with MPK
- **RQ #3:** Does MorphOS maintain **performance?**
 - Throughput

More evaluation in the paper


Experimental Setup and Variants

Setup:


- Intel Xeon 5317, 256GB RAM
- 10G NIC: Intel X520
- Qemu 8.2.6 + VPP 24.06

Variants:

- **Linux + Click:** baseline
- **Unikraft + Click:** native VNF programs
- **MorphOS + MorphClick:** eBPF VNF programs




RQ#1: Lightweight Reconfigurability



MorphOS reduces reconfiguration time and preserves state

RQ#2: Safety Hardening

32 CVEs in eBPF verifiers




MorphOS safety hardening protects against verifier vulnerabilities

RQ#3: Performance — Throughput

Unikernels are
2.9x faster

eBPF is 18%
faster than
parameters

eBPF is 10%
slower with
map lookups



↑ Higher is better

MorphOS is 3% faster than Unikraft without hardening and 3% slower with MPK

Conclusion



How do we achieve safe extensibility of networked unikernel applications?

MorphOS

- Introduces live-reconfigurability with verified eBPF to unikernels
- Provides strong correctness and safety guarantees
 - Verification, MPK hardening



[TUM-DSE/
MorphOS](#)


Impact

- Runtime-extensibility for unikernels
- Broaden the applicability of unikernels to real-world systems

Backup

Memory Safety with MPK

- Separate to different pages
 - eBPF program, stack, maps
 - Packet buffers
- Restrict memory permissions based on current context
- Selectively transport packet buffers into eBPF domain



MPK scalability and probabilistic isolation


- MPK is limited to 16 protection domains
 - Exhausted with per-program and per-buffer domains
 - Existing eBPF hardening with MPK ignores IO buffers
- Naive approach: MPKey virtualization
 - libMPK: fall back to paging for cold set of domains
 - VNFs typically have >16 hot domains

-> Probabilistic isolation

- Re-use protection domains
- Accept permission overlap
- Catch cloud tenants that bypass the verifier


Pages
Unikernel
eBPF program #1
eBPF stack protector
eBPF stack
eBPF maps
...
eBPF program #2
Packet buffer #1
...
Packet buffer #m
...
Packet buffer #13

Performance: Firewall



RQ#3: Performance — Latency

LoadGen: MoonGen (sw timestamps at 100kpps, 64B)



MorphOS and Unikraft reduce median latencies by 22% over Linux

- Central buffer management
 - Shared memory pool between app, OS, and eBPF
 - Zero-copy data transfers
 - Hardware-assisted isolation with MPK
- Event model
 - Run-to-completion model
 - Direct callback invocation
 - Optimized for cloud-native single-core VMs
 - Lock-free parallelism
- Network stack bypass
 - Direct driver access like DPDK
 - Preserves interrupt-based processing unlike DPDK

MorphOS combines kernel-stack bypass with the efficiency of unikernels

- Performance
 - Codesign: Compile app into unikernel
 - Changing the app requires recompilation
- Single-address space
 - Single-address space: reduce context switches
 - Lack of isolation
- Isolating extensions
 - Isolation through static verification
 - Verification is error prone